

# Recommendations\_with\_IBM

November 1, 2019

## 1 Recommendations with IBM

" In this notebook, you will be putting your recommendation skills to use on real data from the IBM Watson Studio platform.

You may either submit your notebook through the workspace here, or you may work from your local machine and submit through the next page. Either way assure that your code passes the project [RUBRIC](#). **Please save regularly.**

By following the table of contents, you will build out a number of different methods for making recommendations that can be used for different situations.

### 1.1 Table of Contents

I. Section ?? II. Section ?? III. Section ?? IV. Section ?? V. Section ?? VI. Section ??

At the end of the notebook, you will find directions for how to submit your work. Let's get started by importing the necessary libraries and reading in the data.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import project_tests as t
import pickle

%matplotlib inline

df = pd.read_csv('data/user-item-interactions.csv')
df_content = pd.read_csv('data/articles_community.csv')
del df['Unnamed: 0']
del df_content['Unnamed: 0']

# Show df to get an idea of the data
df.head()
```

```
[1]:  article_id  title \
0      1430.0  using pixiedust for fast, flexible, and easier...
1      1314.0  healthcare python streaming application demo
2      1429.0  use deep learning for image classification
3      1338.0  ml optimization using cognitive assistant
4      1276.0  deploy your python model as a restful api
```

```

                                email
0  ef5f11f77ba020cd36e1105a00ab868bbdbf7fe7
1  083cbdfa93c8444beaa4c5f5e0f5f9198e4f9e0b
2  b96a4f2e92d8572034b1e9b28f9ac673765cd074
3  06485706b34a5c9bf2a0ecdac41daf7e7654ceb7
4  f01220c46fc92c6e6b161b1849de11faacd7ccb2

```

```
[2]: # Show df_content to get an idea of the data
df_content.head()
```

```
[2]:                                doc_body \
0  Skip navigation Sign in SearchLoading...\r\n\r...
1  No Free Hunch Navigation * kaggle.com\r\n\r\n ...
2  * Login\r\n * Sign Up\r\n\r\n * Learning Pat...
3  DATALAYER: HIGH THROUGHPUT, LOW LATENCY AT SCA...
4  Skip navigation Sign in SearchLoading...\r\n\r...
```

```

                                doc_description \
0  Detect bad readings in real time using Python ...
1  See the forest, see the trees. Here lies the c...
2  Heres this weeks news in Data Science and Bi...
3  Learn how distributed DBs solve the problem of...
4  This video demonstrates the power of IBM DataS...

```

	doc_full_name	doc_status	article_id
0	Detect Malfunctioning IoT Sensors with Streami...	Live	0
1	Communicating data science: A guide to present...	Live	1
2	This Week in Data Science (April 18, 2017)	Live	2
3	DataLayer Conference: Boost the performance of...	Live	3
4	Analyze NY Restaurant data using Spark in DSX	Live	4

### 1.1.1 Part I : Exploratory Data Analysis

Use the dictionary and cells below to provide some insight into the descriptive statistics of the data.

1. What is the distribution of how many articles a user interacts with in the dataset? Provide a visual and descriptive statistics to assist with giving a look at the number of times each user interacts with an article.

```
[3]: # Descriptive statistics
df["email"].value_counts().describe()
```

```
[3]: count    5148.000000
mean         8.930847
std          16.802267
min           1.000000
25%           1.000000
50%           3.000000
```

```

75%          9.000000
max          364.000000
Name: email, dtype: float64

```

```


[4]: # Plot the number of times each user interacts with an article
plt.hist(df["email"].value_counts(), bins=100)
plt.ylabel("# articles")
plt.xlabel("# users")

```

```

[4]: Text(0.5, 0, '# users')

```



```

[5]: # Fill in the median and maximum number of user_article interactions below

median_val = 3 # 50% of individuals interact with ____ number of articles or
→fewer.
max_views_by_user = 364 # The maximum number of user-article interactions by any
→1 user is _____.

```

2. Explore and remove duplicate articles from the **df\_content** dataframe.

```

[6]: # Find and explore duplicate articles
df_content[df_content["article_id"].duplicated()]

```

```

[6]: doc_body \
365 Follow Sign in / Sign up Home About Insight Da...
692 Homepage Follow Sign in / Sign up Homepage * H...
761 Homepage Follow Sign in Get started Homepage *...
970 This video shows you how to construct queries ...

```

```
971 Homepage Follow Sign in Get started * Home\r\n...
```

```
doc_description \
365 During the seven-week Insight Data Engineering...
692 One of the earliest documented catalogs was co...
761 Todays world of data science leverages data f...
970 This video shows you how to construct queries ...
971 If you are like most data scientists, you are ...
```

	doc_full_name	doc_status	article_id
365	Graph-based machine learning	Live	50
692	How smart catalogs can turn the big data flood...	Live	221
761	Using Apache Spark as a parallel processing fr...	Live	398
970	Use the Primary Index	Live	577
971	Self-service data preparation with IBM Data Re...	Live	232

```
[7]: # Remove any rows that have the same article_id - only keep the first
df_content.drop_duplicates(subset="article_id", inplace = True)
```

3. Use the cells below to find:

- a. The number of unique articles that have an interaction with a user.
- b. The number of unique articles in the dataset (whether they have any interactions or not).
- c. The number of unique users in the dataset. (excluding null values)
- d. The number of user-article interactions in the dataset.

```
[8]: unique_articles = df["article_id"].nunique() # The number of unique articles
      ↳ that have at least one interaction
unique_articles
```

```
[8]: 714
```

```
[9]: total_articles = df_content.shape[0] # The number of unique articles on the IBM
      ↳ platform
total_articles
```

```
[9]: 1051
```

```
[10]: unique_users = df["email"].nunique() # The number of unique users
unique_users
```

```
[10]: 5148
```

```
[11]: user_article_interactions = df.shape[0] # The number of user-article interactions
user_article_interactions
```

```
[11]: 45993
```

4. Use the cells below to find the most viewed **article\_id**, as well as how often it was viewed. After talking to the company leaders, the `email_mapper` function was deemed a reasonable way to map users to ids. There were a small number of null values, and it was found that all of these null values likely belonged to a single user (which is how they are stored using the function below).

```
[12]: df["article_id"].value_counts().head()
```

```
[12]: 1429.0    937
      1330.0    927
      1431.0    671
      1427.0    643
      1364.0    627
      Name: article_id, dtype: int64
```

```
[13]: most_viewed_article_id = "1429.0" # The most viewed article in the dataset as a
      ↳ string with one value following the decimal
      max_views = 937 # The most viewed article in the dataset was viewed how many
      ↳ times?
```

```
[14]: ## No need to change the code here - this will be helpful for later parts of the
      ↳ notebook
      # Run this cell to map the user email to a user_id column and remove the email
      ↳ column

      def email_mapper():
          coded_dict = dict()
          cter = 1
          email_encoded = []

          for val in df['email']:
              if val not in coded_dict:
                  coded_dict[val] = cter
                  cter+=1

              email_encoded.append(coded_dict[val])
          return email_encoded

      email_encoded = email_mapper()
      del df['email']
      df['user_id'] = email_encoded

      # show header
      df.head()
```

```
[14]: article_id      title  user_id
0      1430.0  using pixiedust for fast, flexible, and easier...      1
1      1314.0      healthcare python streaming application demo      2
2      1429.0      use deep learning for image classification      3
3      1338.0      ml optimization using cognitive assistant      4
4      1276.0      deploy your python model as a restful api      5
```

```
[15]: ## If you stored all your results in the variable names above,
      ## you shouldn't need to change anything in this cell

      sol_1_dict = {
          '50% of individuals have ____ or fewer interactions.': median_val,
```

```

    'The total number of user-article interactions in the dataset is _____.':
    →user_article_interactions,
    'The maximum number of user-article interactions by any 1 user is _____.':
    → max_views_by_user,
    'The most viewed article in the dataset was viewed ____ times.':
    →max_views,
    'The article_id of the most viewed article is _____.':
    →most_viewed_article_id,
    'The number of unique articles that have at least 1 rating _____.':
    →unique_articles,
    'The number of unique users in the dataset is ____': unique_users,
    'The number of unique articles on the IBM platform': total_articles
}

# Test your dictionary against the solution
t.sol_1_test(sol_1_dict)

```

It looks like you have everything right here! Nice job!

### 1.1.2 Part II: Rank-Based Recommendations

Unlike in the earlier lessons, we don't actually have ratings for whether a user liked an article or not. We only know that a user has interacted with an article. In these cases, the popularity of an article can really only be based on how often an article was interacted with.

1. Fill in the function below to return the **n** top articles ordered with most interactions as the top. Test your function using the tests below.

```

[296]: def get_top_articles(n, df=df):
    """
    INPUT:
    n - (int) the number of top articles to return
    df - (pandas dataframe) df as defined at the top of the notebook

    OUTPUT:
    top_articles - (list) A list of the top 'n' article titles

    """
    top_ids = df["article_id"].value_counts().head(n).index.tolist()
    top_articles = df.set_index("article_id").loc[top_ids, :].
    →reset_index(inplace=False)
    top_articles = list(top_articles["title"].unique())
    return top_articles # Return the top article titles from df (not df_content)

def get_top_article_ids(n, df=df):
    """
    INPUT:
    n - (int) the number of top articles to return
    df - (pandas dataframe) df as defined at the top of the notebook

```

*OUTPUT:*

*top\_articles - (list) A list of the top 'n' article titles*

*'''*

*# Your code here*

```
top_articles = df["article_id"].value_counts().head(n).index.astype(str).  
→tolist()
```

*return top\_articles # Return the top article ids*

```
[297]: print(get_top_articles(10))  
print(get_top_article_ids(10))
```

```
['use deep learning for image classification', 'insights from new york car  
accident reports', 'visualize car data with brunel', 'use xgboost, scikit-learn  
& ibm watson machine learning apis', 'predicting churn with the spss random tree  
algorithm', 'healthcare python streaming application demo', 'finding optimal  
locations of new store using decision optimization', 'apache spark lab, part 1:  
basic concepts', 'analyze energy consumption in buildings', 'gosales  
transactions for logistic regression model']  
['1429.0', '1330.0', '1431.0', '1427.0', '1364.0', '1314.0', '1293.0', '1170.0',  
'1162.0', '1304.0']
```

```
[18]: # Test your function by returning the top 5, 10, and 20 articles  
top_5 = get_top_articles(5)  
top_10 = get_top_articles(10)  
top_20 = get_top_articles(20)  
  
# Test each of your three lists from above  
t.sol_2_test(get_top_articles)
```

Your top\_5 looks like the solution list! Nice job.  
Your top\_10 looks like the solution list! Nice job.  
Your top\_20 looks like the solution list! Nice job.

### 1.1.3 Part III: User-User Based Collaborative Filtering

1. Use the function below to reformat the **df** dataframe to be shaped with users as the rows and articles as the columns.

- Each **user** should only appear in each **row** once.
- Each **article** should only show up in one **column**.
- If a user has interacted with an article, then place a 1 where the user-row meets for that article-column. It does not matter how many times a user has interacted with the article, all entries where a user has interacted with an article should be a 1.

- If a user has not interacted with an item, then place a zero where the user-row meets for that article-column.

Use the tests to make sure the basic structure of your matrix matches what is expected by the solution.

```
[19]: dummies = pd.get_dummies(df["article_id"])
user_item = pd.concat([df, dummies], axis=1)
user_item = user_item.drop([user_item.columns[0], user_item.columns[1]], axis=1)
user_item = user_item.groupby("user_id", as_index=False).sum()
user_item = user_item.set_index("user_id")
```

```
[20]: # create the user-article matrix with 1's and 0's

def create_user_item_matrix(df):
    '''
    INPUT:
    df - pandas dataframe with article_id, title, user_id columns

    OUTPUT:
    user_item - user item matrix

    Description:
    Return a matrix with user ids as rows and article ids on the columns with 1_
    →values where a user interacted with
    an article and a 0 otherwise
    '''
    # Fill in the function here
    user_item = df.groupby(["user_id", "article_id"])["title"].max().unstack()
    user_item = user_item.notnull().astype(int)

    return user_item # return the user_item matrix

user_item = create_user_item_matrix(df)
user_item.head()
```

```
[20]: article_id  0.0      2.0      4.0      8.0      9.0      12.0      14.0      15.0  \
user_id
1           0         0         0         0         0         0         0         0
2           0         0         0         0         0         0         0         0
3           0         0         0         0         0         1         0         0
4           0         0         0         0         0         0         0         0
5           0         0         0         0         0         0         0         0

article_id  16.0      18.0      ...  1434.0  1435.0  1436.0  1437.0  1439.0  \
user_id      ...
1           0         0      ...         0         0         1         0         1
2           0         0      ...         0         0         0         0         0
3           0         0      ...         0         0         1         0         0
4           0         0      ...         0         0         0         0         0
```



5	0	0	...	0	0	0	0	0
---	---	---	-----	---	---	---	---	---

article_id	1440.0	1441.0	1442.0	1443.0	1444.0
user_id					
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0
5	0	0	0	0	0

[5 rows x 714 columns]

```
[21]: ## Tests: You should just need to run this cell. Don't change the code.
assert user_item.shape[0] == 5149, "Oops! The number of users in the_
→user-article matrix doesn't look right."
assert user_item.shape[1] == 714, "Oops! The number of articles in the_
→user-article matrix doesn't look right."
assert user_item.sum(axis=1)[1] == 36, "Oops! The number of articles seen by_
→user 1 doesn't look right."
print("You have passed our quick tests! Please proceed!")
```

You have passed our quick tests! Please proceed!

2. Complete the function below which should take a `user_id` and provide an ordered list of the most similar users to that user (from most similar to least similar). The returned result should not contain the provided `user_id`, as we know that each user is similar to him/herself. Because the results for each user here are binary, it (perhaps) makes sense to compute similarity as the dot product of two users.

Use the tests to test your function.

```
[28]: def find_similar_users(user_id, user_item=user_item):
    '''
    INPUT:
    user_id - (int) a user_id
    user_item - (pandas dataframe) matrix of users by articles:
                1's when a user has interacted with an article, 0 otherwise

    OUTPUT:
    similar_users - (list) an ordered list where the closest users (largest dot_
→product users)
                    are listed first

    Description:
    Computes the similarity of every pair of users based on the dot product
    Returns an ordered

    '''
    # compute similarity of each user to the provided user
```

```

    similarity_vector = user_item[user_item.index == user_id].dot(np.
→transpose(user_item))
    # sort by similarity
    similarity_vector_sorted = similarity_vector.sort_values(axis=1, by=user_id,
→ascending=False)
    # create list of just the ids
    most_similar_users = similarity_vector_sorted.columns.tolist()
    # remove the own user's id
    most_similar_users.remove(user_id)
    return most_similar_users # return a list of the users in order from most to
→least similar

```

```

[29]: # Do a spot check of your function
print("The 10 most similar users to user 1 are: {}".
→format(find_similar_users(1)[:10]))
print("The 5 most similar users to user 3933 are: {}".
→format(find_similar_users(3933)[:5]))
print("The 3 most similar users to user 46 are: {}".
→format(find_similar_users(46)[:3]))

```

The 10 most similar users to user 1 are: [3933, 23, 3782, 203, 4459, 3870, 131, 4201, 46, 5041]

The 5 most similar users to user 3933 are: [1, 23, 3782, 203, 4459]

The 3 most similar users to user 46 are: [4201, 3782, 23]

3. Now that you have a function that provides the most similar users to each user, you will want to use these users to find articles you can recommend. Complete the functions below to return the articles you would recommend to each user.

```

[104]: def get_article_names(article_ids, df=df):
    """
    INPUT:
    article_ids - (list) a list of article ids
    df - (pandas dataframe) df as defined at the top of the notebook

    OUTPUT:
    article_names - (list) a list of article names associated with the list of
→article ids
                    (this is identified by the title column)
    """
    # Your code here
    df_selected = df.loc[df["article_id"].isin(article_ids)]
    article_names = df_selected["title"].unique().tolist()

    return article_names # Return the article names associated with list of
→article ids

```

```

def get_user_articles(user_id, =user_item):
    '''
    INPUT:
    user_id - (int) a user id
    user_item - (pandas dataframe) matrix of users by articles:
                1's when a user has interacted with an article, 0 otherwise

    OUTPUT:
    article_ids - (list) a list of the article ids seen by the user
    article_names - (list) a list of article names associated with the list of
→article ids
                                (this is identified by the doc_full_name column in
→df_content)

    Description:
    Provides a list of the article_ids and article titles that have been seen by
→a user
    '''
    # Your code here
    article_ids = user_item.loc[user_id]
    article_ids = article_ids[article_ids == 1].index.tolist()
    for item in range(len(article_ids)):
        article_ids[item] = str(article_ids[item])
    article_names = get_article_names(article_ids)
    return article_ids, article_names # return the ids and names

def user_user_recs(user_id, m=10):
    '''
    INPUT:
    user_id - (int) a user id
    m - (int) the number of recommendations you want for the user

    OUTPUT:
    recs - (list) a list of recommendations for the user

    Description:
    Loops through the users based on closeness to the input user_id
    For each user - finds articles the user hasn't seen before and provides them
→as recs
    Does this until m recommendations are found

    Notes:
    Users who are the same closeness are chosen arbitrarily as the 'next' user

    For the user where the number of recommended articles starts below m

```

*and ends exceeding m, the last items are chosen arbitrarily*

```
'''
# Your code here
user_ids = find_similar_users(user_id)
recs = df[df["user_id"].isin(user_ids)]["article_id"]
recs = list(set(recs))
recs = recs[:m]
return recs # return your recommendations for this user_id
```

```
[105]: # Check Results
get_article_names(user_user_recs(1, 10)) # Return 10 recommendations for user 1
```

```
[105]: ['got zip code data? prep it for analytics.  ibm watson data lab  medium',
'timeseries data analysis of iot events by using jupyter notebook',
'the greatest public datasets for ai  startup grind',
'3992    using apache spark to predict attack vectors a...\nName: title, dtype:
object',
'detect malfunctioning iot sensors with streaming analytics',
'this week in data science (april 18, 2017)',
'higher-order logistic regression for large datasets',
'apache spark 2.0: extend structured streaming for spark ml',
'data science bowl 2017',
'analyze ny restaurant data using spark in dsx']
```

```
[108]: # Test your functions here - No need to change this code - just run this cell
assert set(get_article_names(['1024.0', '1176.0', '1305.0', '1314.0', '1422.0',
→'1427.0'])) == set(['using deep learning to reconstruct high-resolution
→audio', 'build a python app on the streaming analytics service', 'gosales
→transactions for naive bayes model', 'healthcare python streaming application
→demo', 'use r dataframes & ibm watson natural language understanding', 'use
→xgboost, scikit-learn & ibm watson machine learning apis']), "Oops! Your the
→get_article_names function doesn't work quite how we expect."
assert set(get_article_names(['1320.0', '232.0', '844.0'])) == set(['housing
→(2015): united states demographic measures', 'self-service data preparation
→with ibm data refinery', 'use the cloudant-spark connector in python
→notebook']), "Oops! Your the get_article_names function doesn't work quite how
→we expect."
assert set(get_user_articles(20)[0]) == set(['1320.0', '232.0', '844.0'])
assert set(get_user_articles(20)[1]) == set(['housing (2015): united states
→demographic measures', 'self-service data preparation with ibm data
→refinery', 'use the cloudant-spark connector in python notebook'])
assert set(get_user_articles(2)[0]) == set(['1024.0', '1176.0', '1305.0', '1314.
→0', '1422.0', '1427.0'])
```

```

assert set(get_user_articles(2)[1]) == set(['using deep learning to reconstruct_
→high-resolution audio', 'build a python app on the streaming analytics_
→service', 'gosales transactions for naive bayes model', 'healthcare python_
→streaming application demo', 'use r dataframes & ibm watson natural language_
→understanding', 'use xgboost, scikit-learn & ibm watson machine learning_
→apis'])
print("If this is all you see, you passed all of our tests! Nice job!")

```

If this is all you see, you passed all of our tests! Nice job!

4. Now we are going to improve the consistency of the **user\_user\_recs** function from above.

- Instead of arbitrarily choosing when we obtain users who are all the same closeness to a given user - choose the users that have the most total article interactions before choosing those with fewer article interactions.
- Instead of arbitrarily choosing articles from the user where the number of recommended articles starts below m and ends exceeding m, choose articles with the articles with the most total interactions before choosing those with fewer total interactions. This ranking should be what would be obtained from the **top\_articles** function you wrote earlier.

```

[279]: def get_top_sorted_users(user_id, df=df, user_item=user_item):
        '''
        INPUT:
        user_id - (int)
        df - (pandas dataframe) df as defined at the top of the notebook
        user_item - (pandas dataframe) matrix of users by articles:
                    1's when a user has interacted with an article, 0 otherwise

        OUTPUT:
        neighbors_df - (pandas dataframe) a dataframe with:
                        neighbor_id - is a neighbor user_id
                        similarity - measure of the similarity of each user to the_
→provided user_id
                        num_interactions - the number of articles viewed by the user_
→if a user has viewed an item

        Other Details - sort the neighbors_df by the similarity and then by number_
→of interactions where
                        highest of each is higher in the dataframe

        '''
        # Create empty dataframe
        neighbor_id = []
        similarity = []
        num_interactions = []
        neighbors_df = pd.DataFrame({

```

```

        "neighbor_id": neighbor_id,
        "similarity": similarity,
        "num_interactions": num_interactions
    })

    for item in user_item.index:
        if item != user_id:
            neighbor_id = item
            # Uses similarity from the find_similar_users function above
            similarity = user_item[user_item.index == user_id].dot(np.
→transpose(user_item.loc[item])).values[0]
            num_interactions = user_item.loc[item].values.sum()
            neighbors_df.loc[neighbor_id] =
→[neighbor_id, similarity, num_interactions]

    neighbors_df = neighbors_df.sort_values(by = ["similarity", "neighbor_id"],
→ascending = [False, True])
    return neighbors_df # Return the dataframe specified in the doc_string

def user_user_recs_part2(user_id, m=10):
    """
    INPUT:
    user_id - (int) a user id
    m - (int) the number of recommendations you want for the user

    OUTPUT:
    recs - (list) a list of recommendations for the user by article id
    rec_names - (list) a list of recommendations for the user by article title

    Description:
    Loops through the users based on closeness to the input user_id
    For each user - finds articles the user hasn't seen before and provides them
→as recs
    Does this until m recommendations are found

    Notes:
    * Choose the users that have the most total article interactions
    before choosing those with fewer article interactions.

    * Choose articles with the articles with the most total interactions
    before choosing those with fewer total interactions.

    """
    # Your code here
    neighbors_df = get_top_sorted_users(user_id)
    user_article_ids, user_article_names = get_user_articles(user_id)

```

```

recs = []

top_m_neighbor_ids = neighbors_df["neighbor_id"].head(m).tolist()
for neighbor in top_m_neighbor_ids:
    neighbor_article_ids, neighbor_article_names = get_user_articles(neighbor)
    recs.extend(neighbor_article_ids)

recs = list(set(recs) - set(user_article_ids))
recs = list(set(recs[:m]))
rec_names = get_article_names(recs)

return recs, rec_names

```

```

[280]: # Quick spot check - don't change this code - just use it to test your functions
rec_ids, rec_names = user_user_recs_part2(20, 10)
print("The top 10 recommendations for user 20 are the following article ids:")
print(rec_ids)
print()
print("The top 10 recommendations for user 20 are the following article names:")
print(rec_names)

```

The top 10 recommendations for user 20 are the following article ids:  
['858.0', '1184.0', '1153.0', '681.0', '1014.0', '1432.0', '112.0', '607.0',  
'1278.0', '1424.0']

The top 10 recommendations for user 20 are the following article names:  
['visualize data with the matplotlib library', 'city population by sex, city and  
city type', '1448 i ranked every intro to data science course on...\nName:  
title, dtype: object', 'real-time sentiment analysis of twitter hashtags with  
spark (+ pixiedust)', 'use spark for python to load data and run sql queries',  
'recommendation system algorithms stats and bots', 'airbnb data for analytics:  
vienna calendar', 'analyze starcraft ii replays with jupyter notebooks',  
'building custom machine learning algorithms with apache systemml', 'develop a  
scala spark model on chicago building violations']

5. Use your functions from above to correctly fill in the solutions to the dictionary below. Then test your dictionary against the solution. Provide the code you need to answer each following the comments below.

```

[285]: ### Tests with a dictionary of results
user1_most_sim = get_top_sorted_users(1).iloc[0].neighbor_id # Find the user 1
    ↳ that is most similar to user 1
user131_10th_sim = get_top_sorted_users(131).iloc[9].neighbor_id # Find the 10th 131
    ↳ most similar user to user 131

```

```

[286]: ## Dictionary Test Here
sol_5_dict = {

```

```

    'The user that is most similar to user 1.': user1_most_sim,
    'The user that is the 10th most similar to user 131': user131_10th_sim,
}

t.sol_5_test(sol_5_dict)

```

This all looks good! Nice job!

6. If we were given a new user, which of the above functions would you be able to use to make recommendations? Explain. Can you think of a better way we might make recommendations? Use the cell below to explain a better method for new users.

**Provide your response here.**

7. Using your existing functions, provide the top 10 recommended articles you would provide for the a new user below. You can test your function against our thoughts to make sure we are all on the same page with how we might make a recommendation.

```

[298]: new_user = '0.0'

# What would your recommendations be for this new user '0.0'? As a new user,
→they have no observed articles.
# Provide a list of the top 10 article ids you would give to
new_user_recs = get_top_article_ids(10) # Your recommendations here

[299]: assert set(new_user_recs) == set(['1314.0', '1429.0', '1293.0', '1427.0', '1162.
→0', '1364.0', '1304.0', '1170.0', '1431.0', '1330.0']), "Oops! It makes sense that
→in this case we would want to recommend the most popular articles, because we
→don't know anything about these users."

print("That's right! Nice job!")

```

That's right! Nice job!

#### 1.1.4 Part IV: Content Based Recommendations (EXTRA - NOT REQUIRED)

Another method we might use to make recommendations is to perform a ranking of the highest ranked articles associated with some term. You might consider content to be the **doc\_body**, **doc\_description**, or **doc\_full\_name**. There isn't one way to create a content based recommendation, especially considering that each of these columns hold content related information.

1. Use the function body below to create a content based recommender. Since there isn't one right answer for this recommendation tactic, no test functions are provided. Feel free to change the function inputs if you decide you want to try a method that requires more input values. The input values are currently set with one idea in mind that you may use to make content based recommendations. One additional idea is that you might want to choose the most popular recommendations that meet your 'content criteria', but again, there is a lot of flexibility in how you might make these recommendations.



**1.1.5 This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.**

```
[ ]: def make_content_recs():  
    '''  
    INPUT:  
  
    OUTPUT:  
  
    '''
```

2. Now that you have put together your content-based recommendation system, use the cell below to write a summary explaining how your content based recommender works. Do you see any possible improvements that could be made to your function? Is there anything novel about your content based recommender?

**1.1.6 This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.**

**Write an explanation of your content based recommendation system here.**

3. Use your content-recommendation system to make recommendations for the below scenarios based on the comments. Again no tests are provided here, because there isn't one right answer that could be used to find these content based recommendations.

**1.1.7 This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.**

```
[ ]: # make recommendations for a brand new user  
  
# make a recommendations for a user who only has interacted with article id_  
→ '1427.0'
```

## 1.1.8 Part V: Matrix Factorization

In this part of the notebook, you will build use matrix factorization to make article recommendations to the users on the IBM Watson Studio platform.

1. You should have already created a **user\_item** matrix above in **question 1** of **Part III** above. This first question here will just require that you run the cells to get things set up for the rest of **Part V** of the notebook.

```
[308]: # Load the matrix here  
#user_item_matrix = pd.read_pickle('user_item_matrix.p')  
user_item_matrix = user_item
```

```
[309]: # quick look at the matrix  
user_item_matrix.head()
```

```
[309]: article_id  0.0      2.0      4.0      8.0      9.0      12.0      14.0      15.0      \
user_id
1          0          0          0          0          0          0          0          0
2          0          0          0          0          0          0          0          0
3          0          0          0          0          0          1          0          0
4          0          0          0          0          0          0          0          0
5          0          0          0          0          0          0          0          0
```

```
article_id 16.0      18.0      ... 1434.0 1435.0 1436.0 1437.0 1439.0 \
user_id      ...
1          0          0      ...      0          0          1          0          1
2          0          0      ...      0          0          0          0          0
3          0          0      ...      0          0          1          0          0
4          0          0      ...      0          0          0          0          0
5          0          0      ...      0          0          0          0          0
```

```
article_id 1440.0 1441.0 1442.0 1443.0 1444.0
user_id
1          0          0          0          0          0
2          0          0          0          0          0
3          0          0          0          0          0
4          0          0          0          0          0
5          0          0          0          0          0
```

[5 rows x 714 columns]

2. In this situation, you can use Singular Value Decomposition from [numpy](#) on the user-item matrix. Use the cell to perform SVD, and explain why this is different than in the lesson.

```
[310]: # Perform SVD on the User-Item Matrix Here
u, s, vt = np.linalg.svd(user_item_matrix) # use the built in to get the three
→matrices
```

```
[312]: print(u.shape)
print(s.shape)
print(vt.shape)
```

(5149, 5149)

(714,)

(714, 714)

**Provide your response here.** User\_item\_matrix has no missing values, so the SVD can decompose user-item matrix to the three matrices u, s, vt. If there are any missing values, funkSVD must be used.

3. Now for the tricky part, how do we choose the number of latent features to use? Running the below cell, you can see that as the number of latent features increases, we obtain a lower error rate on making predictions for the 1 and 0 values in the user-item matrix. Run the cell below to get an idea of how the accuracy improves as we increase the number of latent features.

```
[311]: num_latent_feats = np.arange(10,700+10,20)
sum_errs = []


for k in num_latent_feats:
    # restructure with k latent features
    s_new, u_new, vt_new = np.diag(s[:k]), u[:, :k], vt[:k, :]

    # take dot product
    user_item_est = np.around(np.dot(np.dot(u_new, s_new), vt_new))

    # compute error for each prediction to actual value
    diffs = np.subtract(user_item_matrix, user_item_est)

    # total errors and keep track of them
    err = np.sum(np.sum(np.abs(diffs)))
    sum_errs.append(err)

plt.plot(num_latent_feats, 1 - np.array(sum_errs)/df.shape[0]);
plt.xlabel('Number of Latent Features');
plt.ylabel('Accuracy');
plt.title('Accuracy vs. Number of Latent Features');
```



4. From the above, we can't really be sure how many features to use, because simply having a

better way to predict the 1's and 0's of the matrix doesn't exactly give us an indication of if we are able to make good recommendations. Instead, we might split our dataset into a training and test set of data, as shown in the cell below.

Use the code from question 3 to understand the impact on accuracy of the training and test sets of data with different numbers of latent features. Using the split below:

- How many users can we make predictions for in the test set?
- How many users are we not able to make predictions for because of the cold start problem?
- How many articles can we make predictions for in the test set?
- How many articles are we not able to make predictions for because of the cold start problem?

```
[316]: df_train = df.head(40000)
df_test = df.tail(5993)

def create_test_and_train_user_item(df_train, df_test):
    '''
    INPUT:
    df_train - training dataframe
    df_test - test dataframe

    OUTPUT:
    user_item_train - a user-item matrix of the training dataframe
                      (unique users for each row and unique articles for each_
→column)
    user_item_test - a user-item matrix of the testing dataframe
                    (unique users for each row and unique articles for each_
→column)
    test_idx - all of the test user ids
    test_arts - all of the test article ids

    '''
    # Your code here
    user_item_train = create_user_item_matrix(df_train)
    user_item_test = create_user_item_matrix(df_test)

    test_idx = user_item_test.index
    test_arts = user_item_test.columns
    return user_item_train, user_item_test, test_idx, test_arts

user_item_train, user_item_test, test_idx, test_arts =
→create_test_and_train_user_item(df_train, df_test)
```

```
[316]: Int64Index([2917, 3024, 3093, 3193, 3527, 3532, 3684, 3740, 3777, 3801,
...
5140, 5141, 5142, 5143, 5144, 5145, 5146, 5147, 5148, 5149],
```

```
dtype='int64', name='user_id', length=682)
```

```
[328]: #How many users can we make predictions for in the test set?:  
len(set(user_item_train.index).intersection(set(user_item_test.index)))
```

```
[328]: 20
```

```
[330]: #How many users in the test set are we not able to make predictions for because_  
→of the cold start problem?  
len(set(user_item_test.index).difference(set(user_item_train.index)))
```

```
[330]: 662
```

```
[331]: #How many articles can we make predictions for in the test set?  
len(set(user_item_train.columns).intersection(set(user_item_test.columns)))
```

```
[331]: 574
```

```
[332]: #How many articles in the test set are we not able to make predictions for_  
→because of the cold start problem?  
len(set(user_item_test.columns).difference(set(user_item_train.columns)))
```

```
[332]: 0
```

```
[341]: # Replace the values in the dictionary below  
a = 662  
b = 574  
c = 20  
d = 0  
  
sol_4_dict = {  
    'How many users can we make predictions for in the test set?': c,  
    'How many users in the test set are we not able to make predictions for_  
→because of the cold start problem?': a,  
    'How many movies can we make predictions for in the test set?': b,  
    'How many movies in the test set are we not able to make predictions for_  
→because of the cold start problem?': d  
}  
t.sol_4_test(sol_4_dict)
```

Awesome job! That's right! All of the test movies are in the training data, but there are only 20 test users that were also in the training set. All of the other users that are in the test set we have no data on. Therefore, we cannot make predictions for these users using SVD.

5. Now use the **user\_item\_train** dataset from above to find U, S, and V transpose using SVD. Then find the subset of rows in the **user\_item\_test** dataset that you can predict using this matrix decomposition with different numbers of latent features to see how many features makes sense to keep based on the accuracy on the test data. This will require combining what was done in questions 2 - 4.

Use the cells below to explore how well SVD works towards making predictions for recommendations on the test data.

```
[344]: # fit SVD on the user_item_train matrix
u_train, s_train, vt_train = np.linalg.svd(user_item_train, full_matrices=False)
→# fit svd similar to above then use the cells below
```

```
[347]: print(u_train.shape)
print(s_train.shape)
print(vt_train.shape)
```

(4487, 714)

(714,)

(714, 714)

```
[345]: # Use these cells to see how well you can use the training
# decomposition to predict on test data
```

```
[358]: u_test = u_train[user_item_train.index.isin(test_idx)]
vt_test = vt_train[user_item_train.columns.isin(test_arts)].T
print(u_test.shape)
print(vt_test.shape)
```

(20, 714)

(714, 574)

```
[359]: test_users = set(user_item_train.index).intersection(set(user_item_test.index))
test_articles = set(user_item_train.columns).intersection(set(user_item_test.
→columns))
```

```
[364]: user_item_test_subset = user_item_test.loc[test_users, test_articles]
```

```
[369]: num_latent_feats = np.arange(10,700+10,20)
sum_errs_train = []
sum_errs_test = []

for k in num_latent_feats:
    # restructure with k latent features
    s_train_new, u_train_new, vt_train_new = np.diag(s_train[:k]), u_train[:, :
→k], vt_train[:, :]
    u_test_new, vt_test_new = u_test[:, :k], vt_test[:, :]
    # take dot product
    user_item_train_est = np.around(np.dot(np.dot(u_train_new, s_train_new),
→vt_train_new))
    user_item_test_est = np.around(np.dot(np.dot(u_test_new, s_train_new),
→vt_test_new))
    user_item_test_est.shape
    # compute error for each prediction to actual value
    diffs_train = np.subtract(user_item_train, user_item_train_est)
```

```


diffs_test = np.subtract(user_item_test_subset, user_item_test_est)

# total errors and keep track of them
sum_errs_train.append(np.sum(np.sum(np.abs(diffs_train))))
sum_errs_test.append(np.sum(np.sum(np.abs(diffs_test))))

err = np.sum(np.sum(np.abs(diffs)))
sum_errs.append(err)

fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(num_latent_feats, 1 - np.array(sum_errs_train)/(user_item_train.
    →shape[0]*user_item_test.shape[1]), label='Train Data');
plt.plot(num_latent_feats, 1 - np.array(sum_errs_test)/(user_item_test.
    →shape[0]*user_item_test.shape[1]), label='Test Data');
plt.xlabel('Number of Latent Features');
plt.ylabel('Accuracy');
ax.legend(loc="lower right")
plt.title('Accuracy vs. Number of Latent Features');

```



6. Use the cell below to comment on the results you found in the previous question. Given the circumstances of your results, discuss what you might do to determine if the recommendations

you make with any of the above recommendation systems are an improvement to how users currently find articles?

The recommendation system build in this project, shows a high accuracy for training and testing data from above 300 latent features, which is a good result. Using accuracy as a model performance is not always enough. With the given data we can only make predictions for 20 users in the testing set, giving us a problem of inbalance between the training and testing data. Thus, accuracy might not be a good measure for the model performance and another indicators have to be calculated (for example f1 score or recall). To test the real performance of the recommendation system, we suggest to run a A/B test. A simple experiment design could divide the users into two groups: a group A using the new recommendation system and a group B being the control group. The results should testify, if the new recommendation system performs well in practise.

## 1.2 Conclusion

Congratulations! You have reached the end of the Recommendations with IBM project!

**Tip:** Once you are satisfied with your work here, check over your report to make sure that it satisfies all the areas of the [rubric](#). You should also probably remove all of the “Tips” like this one so that the presentation is as polished as possible.

## 1.3 Directions to Submit

Before you submit your project, you need to create a .html or .pdf version of this notebook in the workspace here. To do that, run the code cell below. If it worked correctly, you should get a return code of 0, and you should see the generated .html file in the workspace directory (click on the orange Jupyter icon in the upper left).

Alternatively, you can download this report as .html via the **File > Download as** sub-menu, and then manually upload it into the workspace directory by clicking on the orange Jupyter icon in the upper left, then using the Upload button.

Once you’ve done this, you can submit your project by clicking on the “Submit Project” button in the lower right here. This will create and submit a zip file with this .ipynb doc and the .html or .pdf version you created. Congratulations!

```
[1]: from subprocess import call
    call(['python', '-m', 'nbconvert', 'Recommendations_with_IBM.ipynb'])
```

```
[1]: 0
```

```
[ ]:
```